

These lecture notes include some material from Professors Bertossi, Kolaitis, Guagliardo, Vardi, Libkin, Barland, McMahan

Nested Queries

Lecture Handout

Dr Evgenia Ternovska

Associate Professor

Simon Fraser University

Aggregate results in WHERE

The right way

Account			
Number	Branch	CustID	Balance
111	London	1	1330.00
222	London	2	1756.00
333	Edinburgh	1	450.00

Accounts with a **higher balance** than the **average of all accounts**

```
SELECT A.number
FROM Account A
WHERE A.balance > ( SELECT AVG(A1.balance)
                   FROM Account A1 );
```

Answer:

Number
111
222

Aggregate results in WHERE

The wrong way

Accounts with a higher balance than the average of all accounts

```
SELECT  A.number
FROM    Account A
WHERE   A.balance > AVG( SELECT A1.balance
                        FROM   Account A1 );
```

ERROR in PostgreSQL

Aggregate functions can only be used in **SELECT** and **HAVING**

Check in SQLite

Comparisons with subquery results

```
SELECT
...
FROM
...
WHERE
term op ( subquery ) ;
```

Allowed as long as **subquery** returns a **single value**

```
SELECT
...
FROM
...
WHERE
(term1, ..., termn) op ( subquery ) ;
```

Allowed as long as **subquery** returns a **single row** with n columns

The WHERE clause revisited

term := attribute | value

comparison :=

- ▶ (term, ..., term) **op** (term, ..., term)
with **op** ∈ {=, <>, <, >, <=, >=}
- ▶ term **IS** [**NOT**] **NULL**
- ▶ (term, ..., term) **op ANY** (query)
- ▶ (term, ..., term) **op ALL** (query)
- ▶ (term, ..., term) [**NOT**] **IN** (query)
- ▶ **EXISTS** (query)

condition :=

- ▶ comparison
- ▶ condition1 **AND** condition2
- ▶ condition1 **OR** condition2
- ▶ **NOT** condition

Comparisons between tuples

$$(A_1, \dots, A_n) = (B_1, \dots, B_n) \iff A_1 = B_1 \wedge \dots \wedge A_n = B_n$$

$$(A_1, \dots, A_n) <> (B_1, \dots, B_n) \iff A_1 \neq B_1 \vee \dots \vee A_n \neq B_n$$

$$(A_1, A_2, A_3) < (B_1, B_2, B_3) \quad \text{(generalizes to } n \text{ elements)}$$

$$\iff A_1 < B_1 \vee \left(A_1 = B_1 \wedge \left(A_2 < B_2 \vee (A_2 = B_2 \wedge A_3 < B_3) \right) \right)$$

$$(A_1, A_2, A_3) <= (B_1, B_2, B_3) \quad \text{(generalizes to } n \text{ elements)}$$

$$\iff A_1 < B_1 \vee \left(A_1 = B_1 \wedge \left(A_2 < B_2 \vee (A_2 = B_2 \wedge A_3 \leq B_3) \right) \right)$$

ANY

$(\text{term}, \dots, \text{term}) \text{ op ANY (query)}$

True if **there exists** a row \bar{r} in the results of *query* such that $(\text{term}, \dots, \text{term}) \text{ op } \bar{r}$ is true

Examples:

- ▶ $3 < \text{ANY}(\{1, 2, 3\})$ is false
- ▶ $3 < \text{ANY}(\{2, 3, 4\})$ is true
- ▶ What about $3 < \text{ANY}(\{\})$?

ALL

$(\text{term}, \dots, \text{term}) \text{ op ALL (query)}$

True if **for all** rows \bar{r} in the results of *query* $(\text{term}, \dots, \text{term}) \text{ op } \bar{r}$ is true

Examples:

- ▶ $3 < \text{ALL}(\{5, 4, 6\})$ is true
- ▶ $3 < \text{ALL}(\{4, 3, 5\})$ is false
- ▶ What about $3 < \text{ALL}(\{\})$?

Examples with ANY / ALL

ID of customers from London who own an account

```
SELECT C.custid
FROM   customer C
WHERE  C.city = 'London'
      AND C.custid = ANY ( SELECT A.custid
                           FROM   Account A );
```

Customers living in cities without a branch

```
SELECT *
FROM   customer C
WHERE  C.city <> ALL ( SELECT A.branch
                     FROM   Account A );
```

IN / NOT IN

(term, ..., term) **IN** (query)

same as

(term, ..., term) = **ANY** (query)

(term, ..., term) **NOT IN** (query)

same as

(term, ..., term) <> **ALL** (query)

Examples with IN / NOT IN

ID of customers from London who own an account

```
SELECT C.custid
FROM   customer C
WHERE  C.city = 'London'
      AND C.custid IN ( SELECT A.custid
                        FROM   Account A );
```

Customers living in cities without a branch

```
SELECT *
FROM   customer C
WHERE  C.city NOT IN ( SELECT A.branch
                      FROM   Account A );
```

EXISTS

EXISTS (query) is true if the result of query is **non-empty**

(Stupid) Example

Return all the customers if there are some accounts in London

```
SELECT *
FROM   Customer
WHERE  EXISTS ( SELECT 1
                FROM   Account
                WHERE  branch='London' );
```

Correlated subqueries

All nested queries can refer to attributes in the **parent queries**

(Smarter) Example

Return customers who have an account in London

```
SELECT *
FROM Customer C
WHERE EXISTS ( SELECT 1
                FROM Account A
                WHERE branch='London'
                AND A.custid = C.id );
```

parameters = attributes of a subquery that refer to outer queries

Examples with EXISTS / NOT EXISTS

ID of customers from London who own an account

```
SELECT C.custid
FROM customer C
WHERE C.city = 'London'
AND EXISTS ( SELECT *
              FROM Account A
              WHERE A.custid = C.id);
```

Customers living in cities without a branch

```
SELECT *
FROM customer C
WHERE NOT EXISTS ( SELECT *
                   FROM Account A
                   WHERE A.branch = C.city );
```

Scoping

A subquery has

- ▶ a **local scope** (its **FROM** clause)
- ▶ n **outer scopes** (where n is the **level of nesting**)
(these are the **FROM** clauses of the parent queries)

For each **reference** to an attribute

1. Look for a binding in the local scope
2. If no binding is found, look in the **closest** outer scope
3. If no binding is found, look in the next closest outer scope
4. ...
5. If no binding is found, give error

Attribute bindings

```
SELECT *  
FROM table1  
WHERE EXISTS ( SELECT 1  
                FROM table2  
                WHERE A = B );
```

What A , B refer to depends on the attributes in **table1** and **table2**

- ▶ Always give aliases to tables
- ▶ Always prefix the attributes with the tables they refer to

```
SELECT *  
FROM table1 T1  
WHERE EXISTS ( SELECT 1  
                FROM table2 T2  
                WHERE T2.A = T1.B );
```


The FROM clause revisited

FROM table₁ [[**AS**] T_1], ..., table _{n} [[**AS**] T_n]

table :=

- ▶ base-table
- ▶ join-table
- ▶ (query)

join-table :=

- ▶ table **JOIN** table **ON** condition
- ▶ table **NATURAL JOIN** table
- ▶ table **CROSS JOIN** table

Subqueries in FROM

Must always be given a name (in standard SQL)

```
SELECT * FROM ( SELECT * FROM R );
```

ERROR: subquery in FROM must have an alias

Cannot refer to attributes of other tables in the same **FROM** clause

```
SELECT *  
FROM R, ( SELECT * FROM S WHERE S.a=R.a ) S1 ;
```

ERROR: invalid reference to FROM-clause entry for table "r"

Example: Queries with and without HAVING

Branches with a total balance (across accounts) of at least 500

```
SELECT    A.branch
FROM      Account A
GROUP BY  A.branch
HAVING    SUM(A.balance) >= 500 ;
```

Same query without **HAVING**:

```
SELECT subquery.branch
FROM   ( SELECT    A.branch, SUM(A.balance) AS total
        FROM      Account A
        GROUP BY  A.branch ) AS subquery
WHERE  subquery.total >= 500 ;
```

Example: Aggregation on aggregates

Average of the total balances across each customer's accounts

1. Find the total balance across each customer's accounts
2. Take the average of the totals

```
SELECT    AVG(subquery.tot)
FROM      ( SELECT    A.custid, SUM(A.balance) AS tot
            FROM      Account A
            GROUP BY  A.custid ) AS subquery ;
```

Other bits and pieces of SQL

Ordering

ORDER BY $\langle \text{column}_1 \rangle$ [**DESC**], ..., $\langle \text{column}_n \rangle$ [**DESC**]

Sorts the output rows according to the values of column_1

If two rows have the same value for column_1 ,
they are sorted by the values of column_2 and so on ...

- ▶ Default ordering is **ascending** (can be specified with **ASC**)
- ▶ **Descending** ordering is specified by **DESC**

Ordering example (1)

Account

Number	Branch	CustID	Balance
111	London	1	1330.00
222	London	2	1756.00
333	Edinburgh	1	450.00

```
SELECT *  
FROM Account  
ORDER BY custid ASC, balance DESC ;
```

Number	Branch	CustID	Balance
111	London	1	1330.00
333	Edinburgh	1	450.00
222	London	2	1756.00

21 / 24

Ordering example (2)

Account

Number	Branch	CustID	Balance
111	London	1	1330.00
222	London	2	1756.00
333	Edinburgh	1	450.00

```
SELECT *  
FROM Account  
ORDER BY custid DESC, balance ASC ;
```

Number	Branch	CustID	Balance
222	London	2	1756.00
333	Edinburgh	1	450.00
111	London	1	1330.00

22 / 24

Casting

In most dialects of SQL:

CAST(term **AS** ⟨type⟩)

Rounding

CAST (102,4675 **AS NUMERIC** (5,2)) gives 102.47

Useful also to produce values in a specific format

Aggregation

AVG(**CAST**(term **AS NUMERIC**(p, s)))

avoids rounding errors in some systems

23 / 24

Acknowledgements

[1] Database Systems: The Complete Book, 2nd Edition Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom Prentice Hall, 2009

[2] Database System Concepts, Seventh Edition Avi Silberschatz, Henry F. Korth, S. Sudarshan McGraw-Hill, March 2019 www.db-book.com

Additional references and resources used in preparation of this course are listed on the course webpage or mentioned in slides.

24 / 24